

DataEng S24: Data Storage In-class Assignment

This week you'll gain experience with various ways to load data into a database.

Make a copy of this document and use it to record your results. Store a PDF copy of the document in your git repository along with any needed code before submitting for this week.

The data set for this week is US Census data from 2015. The United States conducts a full census of every household every 10 years (last one was in 2020), but much of the detailed census data comes during the intervening years when the Census Bureau conducts its detailed American Community Survey (ACS) of a randomly selected sample of approximately 3.5 million households each year. The resulting data gives a more detailed view of many factors of American life and the composition of households.

[ACS Census Tract Data for 2015 \(part 1\)](#)

[ACS Census Tract Data for 2015 \(part 2\)](#)

Your job is to load the 2015 data set (approximately 74000 rows divided into two parts). You'll configure a postgres DBMS on a new GCP virtual machine, and then load the two parts of the data using multiple loading methods, comparing the cost of each method. Load the two parts of the data separately so that you gain experience with the issues related to incremental loading of data.

Note that the goal here is not to achieve the fastest load times. Instead, your goal should be to observe and understand a variety of data loading methods. If you start to run out of time, then skip to part I (copy_from) to be sure to get experience with what is probably the fastest way to load bulk data into a Postgres server.

Please highlight your responses so that we can more easily see them

A. [MUST] Discussion Question (discuss as a group near the beginning of the week. Note your own response in this space):

Do you have any experience with ingesting bulk data into a DBMS? If yes, then describe your experience, especially what method was used to input the data into the database. If no, then describe how you might ingest daily incremental breadcrumb data for your class project.

Submit: submit your assignment by this Friday at 10pm

B. [MUST] Configure the Database

1. Create a new GCP virtual machine for this week's work (medium size or larger).
2. Follow the steps listed in the [Installing and Configuring a PostgreSQL server](#) instructions. To keep your project work clean, use a new, different vm and delete the vm when finished with this assignment.
3. Also running the following commands on your VM will help to configure the python module "psycopg2" which you will use to connect to your postgres database:

```
sudo apt install python3 python3-dev python3-venv
sudo apt-get install python3-pip
pip3 install psycopg2-binary
```

- Didn't work, had to run the command:

```
sudo apt install python3-psycopg2
```

C. [MUST] Prepare Data for Loading

1. Copy/upload both ACS Census Tract data files to your VM
2. Create a small test sample by running a linux command similar to the following. The small test sample will help you to quickly test any code that you write before running your code with the full dataset.

```
head -1 acs2015_census_tract_data_part1.csv > AL2015_1.csv
head -1 acs2015_census_tract_data_part2.csv > OR2015_2.csv
grep Alabama acs2015_census_tract_data_part1.csv >> AL2015_1.csv
grep Oregon acs2015_census_tract_data_part2.csv >> OR2015_2.csv
```

The first two commands copy the headers to the sample files and the next command appends Alabama and Oregon 2015 data to the sample files. This should produce files with fewer than 1500 records per file. Use these sample files to save time during testing.

3. Write a python program that connects to your postgres database and creates your main census data table. Start with this example code: [load_inserts.py](#). For example:

```
python3 load_inserts.py -d ./AL2015_1.csv -c
python3 load_inserts.py -d ./OR2015_2.csv
```

- Did work, got an error about the relation not existing. Had to run the command:

```
o python3 load_inserts.py -d ./OR2015_2.csv
  --createtable
```

- I wanted a progress bar so ran the command to being a new virtual env to install tqdm:
 - `python3 -m venv myenv`
 - `source myenv/bin/activate`
- Then I had to setup the missing packages:
 - `pip install tqdm`
 - `pip install psycpg2-binary`
 - `sudo apt-get update`
 - `sudo apt-get install libpq-dev`
 - `sudo apt-get install python3-dev`
 - `pip install psycpg2`
 -

D. [MUST] Baseline - Simple INSERT

The tried and true SQL command [INSERT INTO ...](#) is the most basic way to insert data into a SQL database, and often it is the best choice for small amounts of data, production databases and other situations in which you need to maintain performance and ACID properties of the updated table.

The load_inserts.py program shows how to use simple INSERTs to load data into a database. It is possibly the slowest way to load large amounts of data. For me, it takes approximately 1 second for the Oregon sample and nearly 60 seconds to load each part of the acs data.

Take the program and try it with both the Oregon samples and both parts of the ACS data set. Fill in the appropriate information in the table below (see part J).

```
(myenv) wmm2@week-06-datastorage:~$ python3 load_inserts.py -d ./acs2015_census_tract_data_part1.csv --createtable
Table CensusData created.
Loading 36844 rows
Inserting rows: 100%|██████████████████████████████████████████████████████████████████████████████| 36844/36844 [00:05<00:00, 6266.59it/s]
Finished loading. Elapsed time: 5.8953 seconds.
(myenv) wmm2@week-06-datastorage:~$ python3 load_inserts.py -d ./acs2015_census_tract_data_part2.csv
Loading 37157 rows
Inserting rows: 100%|██████████████████████████████████████████████████████████████████████████████| 37157/37157 [00:06<00:00, 5584.63it/s]
Finished loading. Elapsed time: 6.6738 seconds.
(myenv) wmm2@week-06-datastorage:~$ sudo -u postgres psql postgres
psql (15.6 (Debian 15.6-0+deb12u1))
Type "help" for help.

postgres=# SELECT COUNT(*) FROM CensusData;
 count
-----
 74001
(1 row)

postgres=#
```

- Seems like it took $5.89+6.67=12.56$ seconds total for both parts of the data.

E. [MUST] Disabling Indexes and Constraints

You might notice that the `CensusData` table has a Primary Key constraint and an additional index on the state name column. Indexes and constraints are helpful for query performance but these features can slow down load performance.

Modify your load_inserts.py program to delay creation of constraints/indexes until after the data set is loaded. Enter the resulting load time into the results table below. How much does this technique improve load performance?

```
(myenv) wmm2@week-06-datastorage:~$ python3 load_inserts.py -d ./acs2015_census_tract_data_part1.csv --createtable --optimize-loading  
python3 load_inserts.py -d ./acs2015_census_tract_data_part2.csv  
Table CensusData created with optimization: True.  
Loading 36844 rows  
Inserting rows: 100%|███████████████████████████████████████████| 36844/36844 [00:05<00:00, 6882.52it/s]  
Finished loading. Elapsed time: 5.3731 seconds.  
Indexes and constraints added to CensusData.  
Loading 37157 rows  
Inserting rows: 100%|███████████████████████████████████████████| 37157/37157 [00:08<00:00, 4202.38it/s]  
Finished loading. Elapsed time: 8.8587 seconds.  
(myenv) wmm2@week-06-datastorage:~$ sudo -u postgres psql postgres  
psql (15.6 (Debian 15.6-0+deb12u1))  
Type "help" for help.  
  
postgres=# SELECT COUNT(*) FROM CensusData;  
count  
-----  
74001  
(1 row)  
  
postgres=# \q  
(myenv) wmm2@week-06-datastorage:~$
```

- “Meh”. Seems like with updating the code it’s not much faster running at 5.37+8.85=14.22 seconds. I added an in-line command which tacks on the primary key only if the in-line command isn’t present like this:
if not optimize_loading:
 cursor.execute(f"CREATE INDEX idx_{TableName}_State ON {TableName}(State);")
- I gotta say I’m not too impressed. The speed is just as equally the same total, about 12 seconds. It’s almost as if the second batch took longer.

F. [SHOULD] Disabling Autocommit

You might have noticed that the `load_inserts.py` program sets `autocommit=True` on the database connection. This makes loaded data available to DB queries immediately after each insert. But it also triggers transaction-related overhead operations (e.g., write ahead logging). It also allows readers of the database to view an incomplete set of data during the load.

Modify your load_inserts.py program to avoid setting autocommit=True
Enter the resulting load time into the results table below.

```
(myenv) wmm2@week-06-datastorage:~$ python3 load_inserts.py -d ./acs2015_census_tract_data_part1.c  
sv --createtable --use-single-transaction  
Table CensusData created with optimization: False.  
Loading 36844 rows in a single transaction  
Preparing data: 100%|██████████████████████████████████████| 36844/36844 [00:00<00:00, 101308.94it/s]  
Finished loading. Elapsed time: 4.9816 seconds.  
(myenv) wmm2@week-06-datastorage:~$ python3 load_inserts.py -d ./acs2015_census_tract_data_part2.c  
sv --use-single-transaction  
Loading 37157 rows in a single transaction  
Preparing data: 100%|██████████████████████████████████████| 37157/37157 [00:00<00:00, 106981.35it/s]  
Finished loading. Elapsed time: 4.9164 seconds.  
(myenv) wmm2@week-06-datastorage:~$ sudo -u postgres psql postgres  
psql (15.6 (Debian 15.6-0+deb12u1))  
Type "help" for help.  
  
postgres=# SELECT COUNT(*) FROM CensusData;  
count  
-----  
74001  
(1 row)
```

- Doing this and tossing in all the data at a single time took about just as long because it required a lot of time to prepare the data inorder to get inserted all at the same time with a single call. $4.98+4.91=9.89$ seconds. This might be much faster when tossing in even more data but for this little amount of 74k entries it's really not worth it. Seems like with the clean up it took about the same amount of time as just inserting one at a time.

G. [SHOULD] UNLOGGED table

By default, RDBMS tables incur overheads of write-ahead logging (WAL) such that the database outputs extra metadata about each row insert to a log file known as the Transaction Recovery Log (sometimes just called the WAL or “Write-Ahead Log”). The RDBMS uses that WAL data to recover the contents of the table if/when the RDBMS crashes. Crash Recovery is a great feature but it can slow down load performance.

You can avoid this extra WAL overhead by [loading to an UNLOGGED table](#). Modify your `load_inserts.py` program to load data to an UNLOGGED table. Then enhance `load_inserts.py` to use a SQL query to append the staging data to the main `CensusData` table. Then create the needed index and constraint.

- I tried in various ways, I'm not able to get this to work without getting the error:
`psycopg2.errors.DuplicateTable: relation "idx_censusdata_state" already exists`
- I give up!

H. [ASPIRE] Temp Tables and Memory Tuning

Next compare the above approach with loading the data to [a temporary table](#) (and copying from the temporary table to the CensusData table). Which approach works best for you?

The amount of memory used for temporary tables is default configured to only 8MB. Your VM has enough memory to allocate much more memory to temporary tables. Try allocating 256 MB (or more) to temporary tables. So update the [temp_buffers parameter](#) to allow the database to use more memory for your temporary table. Rerun your load experiments. Did it make a difference?

I. [MUST] Built In Facility (copy_from)

The number one rule of bulk loading is to pay attention to the native facilities provided by the DBMS system implementers. DBMS vendors often put great effort into providing purpose-built loading mechanisms that achieve high performance and scalability.

With a simple, one-server Postgres database, that facility is known as COPY, \copy, or for python programmers [copy_from](#). See the last section of [this blog post](#) for an example.

```
(myenv) wmm2@week-06-datastorage:~$ python3 load_inserts.py -d ./acs2015_census_tract_data_part1.csv --createtable --use-single-transaction --copy-from
Table CensusData created with optimization: False.
Starting to load data from ./acs2015_census_tract_data_part1.csv into CensusData using copy_from.
Data loaded using copy_from method from ./acs2015_census_tract_data_part1.csv.
Total elapsed time: 0.6720 seconds.
(myenv) wmm2@week-06-datastorage:~$ python3 load_inserts.py -d ./acs2015_census_tract_data_part2.csv --use-single-transaction --copy-from
Notice: Table CensusData was not recreated; assuming it exists.
Starting to load data from ./acs2015_census_tract_data_part2.csv into CensusData using copy_from.
Data loaded using copy_from method from ./acs2015_census_tract_data_part2.csv.
Total elapsed time: 0.6858 seconds.
(myenv) wmm2@week-06-datastorage:~$ sudo -u postgres psql postgres
psql (15.6 (Debian 15.6-0+deb12u1))
Type "help" for help.

postgres=# SELECT COUNT(*) FROM CensusData;
 count
-----
 74001
(1 row)

postgres=#
```

- Holy-wow-ness! Look at that!
- This thing took about $0.67+0.68=1.35$ seconds, that's an order of magnitude speed improvement! This is amazing! I gotta get this into my project!

J. [MUST] Results

Use this table to present your results. List only the methods that you actually completed. List only load times for the full amount of Census data (not the small test sample of OR data created in part C). For each loading method, load both parts of the Census data and measure/report the load time for each part separately in the corresponding column of the table.

Method	Time to load part1	Time to load part2
D. Simple inserts	5.89	6.67
E. Drop Indexes and Constraints	5.37	8.85
F. Disable Autocommit	4.98	4.91
G. Use UNLOGGED table	N/A	N/A
H. Temp Table with memory tuning	N/A	N/A
I. copy_from	0.67	0.68

K. [SHOULD] Observations

Use this section to record any observations about the various methods/techniques that you used for bulk loading of the USA Census data. Did you learn anything about why various loading approaches produce varying performance results?

- Yeah that copy_from is the way to go. That thing is awesome!